

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True)

```
data =pd.read_csv('/content/Titanic-Dataset.csv')
```

```
print(data.head())
```

```

  PassengerId  Survived  Pclass  \
0            1         0       3
1            2         1       1
2            3         1       3
3            4         1       1
4            5         0       3

      Name               Sex  Age  SibSp  \
0  Braund, Mr. Owen Harris   male  22.0    1
1  Cumings, Mrs. John Bradley (Florence Briggs Th... female  38.0    1
2    Heikkinen, Miss. Laina  female  26.0    0
3  Futrelle, Mrs. Jacques Heath (Lily May Peel)    female  35.0    1
4    Allen, Mr. William Henry   male  35.0    0

   Parch    Ticket   Fare Cabin Embarked
0      0  A/5 21171   7.2500   NaN        S
1      0    PC 17599  71.2833   C85        C
2      0  STON/O2. 3101282   7.9250   NaN        S
3      0   113803   53.1000  C123        S
4      0  373450   8.0500   NaN        S
```

```
print(data.info())
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
 #   Column             Non-Null Count  Dtype
---  ---
 0   PassengerId        891 non-null    int64
 1   Survived           891 non-null    int64
 2   Pclass             891 non-null    int64
 3   Name               891 non-null    object
 4   Sex               891 non-null    object
 5   Age               714 non-null    float64
 6   SibSp             891 non-null    int64
 7   Parch             891 non-null    int64
 8   Ticket            891 non-null    object
 9   Fare             891 non-null    float64
10   Cabin            204 non-null    object
11   Embarked         889 non-null    object
dtypes: float64(2), int64(5), object(5)
memory usage: 83.7+ KB
None
```

```
print(data.describe())
```

```

   PassengerId  Survived  Pclass    Age  SibSp  \
count  891.000000  891.000000  891.000000  714.000000  891.000000
mean    446.000000    0.383838    2.308642   29.699118    0.523008
std    257.353842    0.486592    0.836071   14.526497    1.102743
min      1.000000    0.000000    1.000000    0.420000    0.000000
25%    223.500000    0.000000    2.000000   20.125000    0.000000
50%    446.000000    0.000000    3.000000   28.000000    0.000000
75%    668.500000    1.000000    3.000000   38.000000    1.000000
max    891.000000    1.000000    3.000000   80.000000    8.000000

   Parch    Fare
count  891.000000  891.000000
mean     0.381594   32.204208
std     0.806057   49.693429
min     0.000000    0.000000
25%     0.000000    7.910400
50%     0.000000   14.454200
75%     0.000000   31.000000
max     6.000000  512.329200
```

```
print(data.isnull().sum())
```

```
PassengerId    0
Survived        0
Pclass         0
Name           0
Sex            0
Age           177
SibSp          0
Parch          0
Ticket         0
Fare           0
Cabin         687
Embarked       2
dtype: int64
```

```
data['Age'].fillna(data['Age'].mean(), inplace=True)
```

/tmp/ipython-input-1981374165.py:1: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chain. The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are set

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = c

```
data['Age'].fillna(data['Age'].mean(), inplace=True)
```

```
data['Embarked'].fillna(data['Embarked'].mode()[0], inplace=True)
```

/tmp/ipython-input-4247733614.py:1: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chain. The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are set

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = c

```
data['Embarked'].fillna(data['Embarked'].mode()[0], inplace=True)
```

```
data.drop('Cabin', axis=1, inplace=True)
```

```
data['Sex'] = data['Sex'].map({'male': 0, 'female': 1})
```

```
data=pd.get_dummies(data, columns=['Embarked'], drop_first=True)
```

```
x = data.drop(['Survived', 'Name', 'Ticket'], axis=1)
y = data['Survived']
```

◆ Gemini

```
from sklearn.model_selection import train_test_split

# Drop the 'Ticket' column from x before splitting
x = x.drop('Ticket', axis=1)

x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_state=42)
```

```
from sklearn.linear_model import LogisticRegression
```

```
model = LogisticRegression(max_iter=100)
```

```
model.fit(x_train, y_train) # Removed the redundant fit call
```

/usr/local/lib/python3.12/dist-packages/sklearn/linear_model/_logistic.py:465: ConvergenceWarning: lbfgs failed to converge (stopping criterion not reached). STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

```
    LogisticRegression
    LogisticRegression()
```

```
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
```

```
y_pred = model.predict(x_test)
```

```
print("Accuracy:", accuracy_score(y_test, y_pred))
```

```
Accuracy: 0.8044692737430168
```

```
print("prediction for first 5 passengers:", model.predict(x_test[:5]))
```


```
prediction for first 5 passengers: [0 0 0 1 1]
```

```
print("actual:", y_test[:5].values)
```

```
actual: [1 0 0 1 1]
```

```
cm=confusion_matrix(y_test, y_pred)
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
plt.title('Confusion Matrix')
plt.xlabel('Predicted')

plt.ylabel('Actual')
plt.show()
```






```
print("classification report:\n",classification_report(y_test, y_pred))
```

```
classification report:
              precision    recall  f1-score   support

     0       0.81         0.87         0.84         105
     1       0.79         0.72         0.75          74

 accuracy          0.80
 macro avg         0.80         0.79         0.80         179
weighted avg         0.80         0.80         0.80         179
```

```
sns.countplot(x='Survived',data=data)
plt.title("Survival Count(0=Died,1=Survived)")
sns.countplot(x='Survived',hue='Sex',data=data)
plt.title("Survival Count by Gender")
plt.show()
sns.countplot(x='Survived',hue='Pclass',data=data)
plt.title("Survival Count by Passenger Class")
plt.show()
sns.histplot(data['Age'],bins=30,kde=True)
plt.title("Age Distribution")
plt.show()
```



```
from sklearn.ensemble import RandomForestClassifier
```

```
rf_model = RandomForestClassifier(n_estimators=100, random_state=42)
```

```
rf_model.fit(x_train, y_train)  
rf_pred = rf_model.predict(x_test)
```

```
print(classification_report(y_test, rf_pred))
print(confusion_matrix(y_test, rf_pred))
print("Accuracy:", accuracy_score(y_test, rf_pred))
```

	precision	recall	f1-score	support
0	0.82	0.87	0.84	105
1	0.79	0.73	0.76	74
accuracy			0.81	179
macro avg	0.81	0.80	0.80	179
weighted avg	0.81	0.81	0.81	179

```
[[91 14]
 [20 54]]
Accuracy: 0.8100558659217877
```

```
import joblib
```

```
joblib.dump(rf_model, "titanic_model.pkl")
```

```
['titanic_model.pkl']
```

Start coding or [generate](#) with AI.